

ANN Implementation

```
In [23]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.layers import Dense, Input
```

```
In [2]: # Load the dataset
file_path = 'Clustered_Customer_Data.csv'
data = pd.read_csv(file_path)
```

Step 1: Data Preprocessing

```
In [3]: # Drop the Unnamed column and handle missing values by filling with the median
data = data.drop(columns=['Unnamed: 0'])
data.fillna(data.median(), inplace=True)
```

```
In [4]: # Separate features and target variable
X = data.drop(columns=['Cluster'])
y = data['Cluster']
```

```
In [5]: # Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [21]: print(X_train.shape)
print(y_train.shape)

(40000, 17)
(40000,)
```

Step 2: Data Splitting

```
In [7]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

Step 3: Building the ANN Model

```
In [24]: model = Sequential()

# Input layer (separate)
model.add(Input(shape=(17,)))

# First hidden layer
model.add(Dense(32, activation='relu'))

# Additional hidden layer
model.add(Dense(16, activation='relu'))

# Output layer with units equal to the number of unique clusters, using softmax
model.add(Dense(len(y.unique()), activation='softmax'))

model.summary()
```

Model: "sequential_4"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense_3 (Dense) | (None, 32) | 576 |
| dense_4 (Dense) | (None, 16) | 528 |
| dense_5 (Dense) | (None, 3) | 51 |

Total params: 1,155 (4.51 KB)

Trainable params: 1,155 (4.51 KB)

Non-trainable params: 0 (0.00 B)

```
In [25]: # Step 4: Compiling the Model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
In [26]: # Step 5: Training the Model
history = model.fit(X_train, y_train, epochs=50, batch_size=32, validation_split=0.2, verbose=1)

Epoch 1/50
1000/1000 — 3s 2ms/step - accuracy: 0.8264 - loss: 0.4043 - val_accuracy: 0.9877 - val_loss: 0.0374
Epoch 2/50
1000/1000 — 2s 2ms/step - accuracy: 0.9876 - loss: 0.0361 - val_accuracy: 0.9891 - val_loss: 0.0260
Epoch 3/50
1000/1000 — 2s 2ms/step - accuracy: 0.9918 - loss: 0.0234 - val_accuracy: 0.9894 - val_loss: 0.0225
Epoch 4/50
1000/1000 — 4s 3ms/step - accuracy: 0.9937 - loss: 0.0184 - val_accuracy: 0.9956 - val_loss: 0.0148
Epoch 5/50
1000/1000 — 2s 2ms/step - accuracy: 0.9952 - loss: 0.0147 - val_accuracy: 0.9967 - val_loss: 0.0118
Epoch 6/50
1000/1000 — 2s 2ms/step - accuracy: 0.9951 - loss: 0.0136 - val_accuracy: 0.9973 - val_loss: 0.0106
Epoch 7/50
1000/1000 — 3s 2ms/step - accuracy: 0.9961 - loss: 0.0114 - val_accuracy: 0.9971 - val_loss: 0.0095
Epoch 8/50
1000/1000 — 3s 2ms/step - accuracy: 0.9965 - loss: 0.0095 - val_accuracy: 0.9965 - val_loss: 0.0090
Epoch 9/50
1000/1000 — 2s 2ms/step - accuracy: 0.9965 - loss: 0.0092 - val_accuracy: 0.9983 - val_loss: 0.0067
Epoch 10/50
1000/1000 — 3s 3ms/step - accuracy: 0.9970 - loss: 0.0087 - val_accuracy: 0.9977 - val_loss: 0.0071
Epoch 11/50
1000/1000 — 4s 2ms/step - accuracy: 0.9970 - loss: 0.0080 - val_accuracy: 0.9959 - val_loss: 0.0122
Epoch 12/50
1000/1000 — 2s 2ms/step - accuracy: 0.9969 - loss: 0.0076 - val_accuracy: 0.9989 - val_loss: 0.0053
Epoch 13/50
1000/1000 — 3s 2ms/step - accuracy: 0.9971 - loss: 0.0068 - val_accuracy: 0.9989 - val_loss: 0.0044
Epoch 14/50
1000/1000 — 2s 2ms/step - accuracy: 0.9982 - loss: 0.0054 - val_accuracy: 0.9970 - val_loss: 0.0059
Epoch 15/50
1000/1000 — 3s 3ms/step - accuracy: 0.9979 - loss: 0.0062 - val_accuracy: 0.9990 - val_loss: 0.0038
Epoch 16/50
1000/1000 — 2s 2ms/step - accuracy: 0.9985 - loss: 0.0047 - val_accuracy: 0.9986 - val_loss: 0.0046
Epoch 17/50
1000/1000 — 3s 2ms/step - accuracy: 0.9980 - loss: 0.0054 - val_accuracy: 0.9992 - val_loss: 0.0031
Epoch 18/50
1000/1000 — 2s 2ms/step - accuracy: 0.9981 - loss: 0.0051 - val_accuracy: 0.9983 - val_loss: 0.0048
Epoch 19/50
1000/1000 — 3s 2ms/step - accuracy: 0.9984 - loss: 0.0045 - val_accuracy: 0.9990 - val_loss: 0.0034
Epoch 20/50
1000/1000 — 3s 3ms/step - accuracy: 0.9985 - loss: 0.0044 - val_accuracy: 0.9984 - val_loss: 0.0033
Epoch 21/50
1000/1000 — 4s 2ms/step - accuracy: 0.9989 - loss: 0.0035 - val_accuracy: 0.9991 - val_loss: 0.0033
Epoch 22/50
1000/1000 — 2s 2ms/step - accuracy: 0.9981 - loss: 0.0051 - val_accuracy: 0.9967 - val_loss: 0.0070
Epoch 23/50
1000/1000 — 3s 2ms/step - accuracy: 0.9986 - loss: 0.0045 - val_accuracy: 0.9995 - val_loss: 0.0025
Epoch 24/50
1000/1000 — 2s 2ms/step - accuracy: 0.9981 - loss: 0.0048 - val_accuracy: 0.9977 - val_loss: 0.0049
Epoch 25/50
1000/1000 — 4s 3ms/step - accuracy: 0.9979 - loss: 0.0051 - val_accuracy: 0.9992 - val_loss: 0.0024
Epoch 26/50
1000/1000 — 4s 2ms/step - accuracy: 0.9989 - loss: 0.0036 - val_accuracy: 0.9992 - val_loss: 0.0022
Epoch 27/50
1000/1000 — 3s 2ms/step - accuracy: 0.9982 - loss: 0.0059 - val_accuracy: 0.9989 - val_loss: 0.0029
Epoch 28/50
1000/1000 — 2s 2ms/step - accuracy: 0.9991 - loss: 0.0032 - val_accuracy: 0.9952 - val_loss: 0.0086
Epoch 29/50
1000/1000 — 4s 3ms/step - accuracy: 0.9980 - loss: 0.0051 - val_accuracy: 0.9990 - val_loss: 0.0025
Epoch 30/50
1000/1000 — 4s 2ms/step - accuracy: 0.9983 - loss: 0.0047 - val_accuracy: 0.9979 - val_loss: 0.0040
Epoch 31/50
1000/1000 — 3s 2ms/step - accuracy: 0.9988 - loss: 0.0035 - val_accuracy: 0.9990 - val_loss: 0.0026
Epoch 32/50
1000/1000 — 2s 2ms/step - accuracy: 0.9993 - loss: 0.0027 - val_accuracy: 0.9969 - val_loss: 0.0080
Epoch 33/50
1000/1000 — 3s 3ms/step - accuracy: 0.9983 - loss: 0.0043 - val_accuracy: 0.9996 - val_loss: 0.0016
Epoch 34/50
1000/1000 — 3s 3ms/step - accuracy: 0.9984 - loss: 0.0040 - val_accuracy: 0.9996 - val_loss: 0.0019
Epoch 35/50
1000/1000 — 2s 2ms/step - accuracy: 0.9985 - loss: 0.0036 - val_accuracy: 0.9998 - val_loss: 0.0019
Epoch 36/50
1000/1000 — 2s 2ms/step - accuracy: 0.9990 - loss: 0.0028 - val_accuracy: 0.9969 - val_loss: 0.0101
Epoch 37/50
1000/1000 — 3s 2ms/step - accuracy: 0.9988 - loss: 0.0043 - val_accuracy: 0.9994 - val_loss: 0.0017
Epoch 38/50
1000/1000 — 2s 2ms/step - accuracy: 0.9992 - loss: 0.0027 - val_accuracy: 0.9992 - val_loss: 0.0023
Epoch 39/50
1000/1000 — 3s 2ms/step - accuracy: 0.9993 - loss: 0.0023 - val_accuracy: 0.9966 - val_loss: 0.0080
Epoch 40/50
1000/1000 — 3s 3ms/step - accuracy: 0.9978 - loss: 0.0057 - val_accuracy: 0.9990 - val_loss: 0.0027
Epoch 41/50
1000/1000 — 2s 2ms/step - accuracy: 0.9993 - loss: 0.0023 - val_accuracy: 0.9992 - val_loss: 0.0021
Epoch 42/50
1000/1000 — 3s 2ms/step - accuracy: 0.9994 - loss: 0.0022 - val_accuracy: 0.9996 - val_loss: 0.0013
Epoch 43/50
1000/1000 — 2s 2ms/step - accuracy: 0.9988 - loss: 0.0033 - val_accuracy: 0.9998 - val_loss: 0.0012
Epoch 44/50
1000/1000 — 3s 2ms/step - accuracy: 0.9989 - loss: 0.0034 - val_accuracy: 0.9996 - val_loss: 0.0017
Epoch 45/50
1000/1000 — 3s 2ms/step - accuracy: 0.9990 - loss: 0.0029 - val_accuracy: 0.9996 - val_loss: 0.0015
Epoch 46/50
1000/1000 — 2s 2ms/step - accuracy: 0.9990 - loss: 0.0028 - val_accuracy: 0.9989 - val_loss: 0.0036
Epoch 47/50
1000/1000 — 2s 2ms/step - accuracy: 0.9987 - loss: 0.0030 - val_accuracy: 0.9995 - val_loss: 0.0016
Epoch 48/50
1000/1000 — 3s 2ms/step - accuracy: 0.9994 - loss: 0.0024 - val_accuracy: 0.9987 - val_loss: 0.0035
Epoch 49/50
1000/1000 — 3s 2ms/step - accuracy: 0.9987 - loss: 0.0036 - val_accuracy: 0.9999 - val_loss: 9.2421e-04
Epoch 50/50
1000/1000 — 3s 3ms/step - accuracy: 0.9993 - loss: 0.0021 - val_accuracy: 0.9992 - val_loss: 0.0022
```

```
In [27]: # Step 6: Evaluating the Model
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"Test Accuracy: {test_accuracy:.2f}")

Test Accuracy: 1.00
```

```
In [28]: # Optional: Predict and display the results on the test set
y_pred = model.predict(X_test)
y_pred_classes = y_pred.argmax(axis=1)

# Directly use y_test as it is already in the correct format
y_test_classes = y_test

313/313 — 0s 1ms/step
```

```
In [29]: # Confusion matrix and classification report
from sklearn.metrics import confusion_matrix, classification_report
print("Confusion Matrix:")
print(confusion_matrix(y_test_classes, y_pred_classes))
print("\nClassification Report:")
print(classification_report(y_test_classes, y_pred_classes))
```

```
Confusion Matrix:
[[1714  0  0]
 [  0 1363  6]
 [  3  3 6911]]

Classification Report:
              precision    recall  f1-score   support

    0           1.00        1.00        1.00        1714
    1           1.00        1.00        1.00        1369
    2           1.00        1.00        1.00         6917

 accuracy          1.00        1.00        1.00       10000
 macro avg          1.00        1.00        1.00       10000
 weighted avg          1.00        1.00        1.00       10000
```

Automating ANN Model Using Keras Tuner

```
In [32]: import pandas as pd
import numpy as np
from tensorflow.keras import Sequential, Input
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras import initializers
import keras_tuner
from keras import *
from keras.layers import *
import tensorflow
from tensorflow import keras
```

Building ANN Architecture

```
In [33]: def build_model(hp):
    model = Sequential()

    # Input layer
    model.add(Input(shape=(17,)))

    # Weight initializer choices
    initializer = hp.Choice("initializer", values=['he_normal', 'glorot_normal', 'random_normal'])

    # First hidden layer
    num_units1 = hp.Choice("num_units_1", values=[50, 100, 150, 200])
    activation1 = hp.Choice("activation_1", values=['relu', 'tanh', 'sigmoid'])
    model.add(Dense(num_units1, activation=activation1,
                    kernel_initializer=getattr(initializers, initializer)(seed=42)))

    # Second hidden layer
    num_units2 = hp.Choice("num_units_2", values=[50, 100, 150, 200])
    activation2 = hp.Choice("activation_2", values=['relu', 'tanh', 'sigmoid'])
    model.add(Dense(num_units2, activation=activation2,
                    kernel_initializer=getattr(initializers, initializer)(seed=42)))

    # Output layer: Len of Unique neuron for Multiclass binary classification
    model.add(Dense(len(y.unique()), activation='sigmoid',
                    kernel_initializer=getattr(initializers, initializer)(seed=42)))

    # Optimizer choices
    opt = hp.Choice("optimizer", values=['adam', 'rmsprop', 'sgd'])

    # Compiling the model
    model.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics=['accuracy'])

    return model
```

Max Trial

4 choices for num_units_1

3 choices for activation_1

4 choices for num_units_2

3 choices for activation_2

3 choices for initializer

3 choices for optimizer

Then the total possible combinations are:

4×3×4×3×3=432

Setting up the tuner

```
In [34]: tuner = keras_tuner.RandomSearch(
    build_model,                # The model-building function with automated hyperparameters
    objective='val_loss',       # Objective to minimize (validation loss)
    max_trials=432,             # Maximum number of trials to run
    directory='TANSeen',        # Directory to save tuning results
    project_name='ANN Market Segmentation' # Name of the project
)
```

Running the tuner

```
In [37]: tuner.search(X_train, y_train, epochs=10, validation_data=(X_test, y_test))

Trial 432 Complete [00h 00m 47s]
val_loss: 0.005716746672987938

Best val_loss So Far: 0.0019498508190736175
Total elapsed time: 04h 46m 57s
```

Retrieving the best hyperparameters

```
In [38]: best_hyperparameters = tuner.get_best_hyperparameters(num_trials=1)[0].values
```

Output the best hyperparameter values

```
In [39]: best_hyperparameters
{'initializer': 'he_normal',
 'num_units_1': 200,
 'activation_1': 'relu',
 'num_units_2': 50,
 'activation_2': 'sigmoid',
 'optimizer': 'adam'}
```

```
In [44]: # Get the best model from the tuner
best_model = tuner.get_best_models(num_models=1)[0]

# Evaluate the best model
loss, accuracy = best_model.evaluate(X_test, y_test, verbose=0)

print(f"Best Model - Validation Loss: {loss:.4f}")
print(f"Best Model - Accuracy: {accuracy:.4f}")

# Print best hyperparameters again for clarity
print("Best Hyperparameters:")
for key, value in best_hyperparameters.items():
    print(f"{key}: {value}")

Best Model - Validation Loss: 0.0019
Best Model - Accuracy: 0.9999
Best Hyperparameters:
initializer: he_normal
num_units_1: 200
activation_1: relu
num_units_2: 50
activation_2: sigmoid
optimizer: adam
```

```
In [45]: # Assuming 'history' is the result of model.fit()
plt.figure(figsize=(12, 4)) # Adjust figure size as needed

plt.subplot(1, 2, 1) # Subplot for accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
```

